

# Model-Free Reinforcement Learning

## Learning Optimal Value Functions

Matthia Sabatelli

November 16, 2021

# Today's Agenda

- ① Model-Free Reinforcement Learning
- ② Monte Carlo Methods
- ③ Temporal Difference Learning

# Recap

## Key Concepts

# Model-Free Reinforcement Learning

So far we have seen:

- How to **estimate** an optimal value function and **discover** an optimal policy  $\pi^*$
- We did this under the assumption that the environment  $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, r \rangle$  was **known**

## Plot Twist!

We will now consider situations where **no complete knowledge** is available:

- Set of possible states  $\mathcal{S}$  ✓
- Set of possible actions  $\mathcal{A}$  ✓
- Transition Function  $\mathcal{P} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$  ✗
- Reward Function  $\mathcal{R} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$  ✗

# Model-Free Reinforcement Learning

When parts of the MDP  $\mathcal{M}$  are **unknown** we do not deal with Dynamic Programming methods anymore but with Reinforcement Learning algorithms!

- We need to overcome the **lack of information** of  $\mathcal{M}$
- We can do this through **experience**
- Gathering experience means **sampling** states  $s$ , actions  $a$  and rewards  $r$  from the environment
- Recall the concept of **trajectory**  $\tau \langle s_t, a_t, r_t, s_{t+1} \rangle$  seen in Lecture 1!

# Model-Free Reinforcement Learning

What does it mean in practice?

- We do not know **the consequences** of our actions
- We do not know the **dynamics** of the environment we are interacting with
- We are no longer **computing** value functions but rather **learning** them

The transition function  $\mathcal{P}$  and the reward function  $\mathcal{R}$  are usually called the **model** of the environment

$\mathcal{P}$  and  $\mathcal{R}$  can be **learned**  $\Rightarrow$  model-based Reinforcement Learning

# Model-Free Reinforcement Learning

Why is model-free Reinforcement Learning so interesting?

- We learn without any **prior knowledge** of the environment
- Trajectories, and therefore **experience**, is sufficient for learning
- We "only" need a value function  $Q(s, a)$ , which is arguably **easier** to learn than the model

The effectiveness of model-free Reinforcement Learning algorithms **highly depends** from how much experience the agent is able to gather!

# Monte Carlo (MC) Methods

Monte Carlo methods:

- Can be used for learning  $V^\pi(s)$  as well as  $Q^\pi(s, a)$
- Although in this lecture we only focus on learning  $V^\pi(s)$
- The key idea is to learn through **sampling returns**

## Assumption!

We assume that we are always dealing with **episodic tasks** i.e. episodes eventually **terminate**.

$$\langle (s_t, a_t, r_t, s_{t+1}) \rangle, t = 0, \dots, T - 1$$



# Monte Carlo (MC) Methods

Let us consider the notion of expected discounted return introduced in Lecture 1:

$$\begin{aligned} G_t &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \\ &= \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}. \end{aligned}$$

- The **goal** is to learn the state-value function of a given policy  $V^{\pi}(s)$ : [Monte Carlo Prediction](#)
- We do this with respect to the  $G_t$  that is obtained by following  $\pi$

# Monte Carlo (MC) Methods

Learning  $V^\pi(s)$  involves the following [steps](#):

- Before learning, each state has its own value  $V(s_t)$
- We follow policy  $\pi$  until an episode terminates
- We compute the discounted return  $G_t$  that was obtained by  $\pi$

## Monte Carlo (MC) Update Rule


We change the value of each state  $V(s_t)$  based on  $G_t$

$$V(s_t) := V(s_t) + \alpha [G_t - V(s_t)].$$

where  $\alpha \in [0, 1]$  is the learning rate parameter.

# Monte Carlo (MC) Methods

Let us consider the following MDP:



- We have policy  $\pi : s_0 \rightarrow s_1 \rightarrow s_0 \rightarrow s_2$
- $\pi$  results in the sequence of rewards  $-1, +2, +3$
- The starting value of each state is 0,  $\gamma = 0.99$  and  $\alpha = 0.5$

# Monte Carlo (MC) Methods

We know that  $G_t$  for starting in  $s_0$  is

$$\begin{aligned} G_t &= \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \\ &= -1 + \gamma 2 + \gamma^2 3 \approx 3.92 \end{aligned}$$

Therefore

$$\begin{aligned} V(s_t) &:= V(s_t) + \alpha [G_t - V(s_t)] \\ V(0) &:= V(0) + \alpha [3.92 - V(0)] \approx 1.96 \end{aligned}$$

# Monte Carlo (MC) Methods

More about Monte Carlo Methods:

- We have only considered the **prediction** case of learning  $V^\pi(s)$
- Monte Carlo **control** algorithms learn an approximation of  $\pi^*$  through  $Q^\pi(s, a)$
- They follow the ideas of **Dynamic Programming** seen in the previous lecture
- The key idea of **sampling real returns** remains!

# Monte Carlo (MC) Methods

Pros & Cons of Monte Carlo Methods:

- Yield **unbiased** updates thanks to  $G_t$  ✓
- **Scale well** to function approximators ✓
- Learning can be very **slow** as one has to wait until the very end of an episode ✗
- There can be **large variance** in the value updates ✗

# Temporal Difference (TD) Learning

*"The simplest and most elegant idea of Reinforcement Learning ..."*

# Temporal Difference (TD) Learning

With TD-Learning methods we **we do not** have to wait until the end of an episode before updating a value estimate

- We only need to wait until the **next step**
- At  $t + 1$  we immediately create a **target** for learning called the TD-target
- We do this by using the observed reward  $r_t$  and the estimate  $V(s_{t+1})$



# Temporal Difference (TD) Learning

Let us again consider the problem of estimating  $V^\pi(s)$ :

## TD Prediction

We change the value of each state  $V(s_t)$  with respect to  $t + 1$  only:

$$V(s_t) := V(s_t) + \alpha[r_t + \gamma V(s_{t+1}) - V(s_t)].$$

- It is clear that we only learn by *looking ahead* in the future one single step
- This is called TD(0) or *one-step TD*

# Temporal Difference (TD) Learning

The key idea of TD-Learning is to learn through **bootstrapping**:

- We update the value of a state with respect to the value of its **successor** state only
- Ideally we would like to use  $V^\pi(s_{t+1})$  for learning but it is unfortunately **unknown**

$$V(s_t) := V(s_t) + \alpha [r_t + \gamma V^\pi(s_{t+1}) - V(s_t)].$$

Therefore we **replace** with a guess instead:

$$V(s_t) := V(s_t) + \alpha [r_t + \gamma V(s_{t+1}) - V(s_t)].$$

# Temporal Difference (TD) Learning

The key idea of TD-Learning is to learn through **bootstrapping**:

- We update the value of a state with respect to the value of its **successor** state only
- Ideally we would like to use  $V^\pi(s_{t+1})$  for learning but it is unfortunately **unknown**


$$V(s_t) := V(s_t) + \alpha [r_t + \gamma V^\pi(s_{t+1}) - V(s_t)].$$

Therefore we **replace** with a guess instead:

$$V(s_t) := V(s_t) + \alpha [r_t + \gamma V(s_{t+1}) - V(s_t)].$$


# Temporal Difference (TD) Learning

Let us go back to the previous MDP:



- We again have policy  $\pi : s_0 \rightarrow s_1 \rightarrow s_0 \rightarrow s_2$
- $\pi$  results in the sequence of rewards  $-1, +2, +3$
- The starting value of each state is 0,  $\gamma = 0.99$  and  $\alpha = 0.5$

# Temporal Difference (TD) Learning



Our first state transition is  $s_0 \rightarrow s_1$

$$V(s_t) := V(s_t) + \alpha[r_t + \gamma V(s_{t+1}) - V(s_t)]$$

$$V(s_0) := V(s_0) + \alpha[r_t + \gamma V(s_1) - V(s_0)]$$

$$V(s_0) := -0.5$$

# Final Slide!

Lecture Takeaway